

Machine Learning

Michaël Clément

`michael.clement@enseirb-matmeca.fr`

2024–25

Bordeaux INP – ENSEIRB-MATMECA – 3A Intelligence Artificielle

Acknowledgements:

- Aurélie Bugeau (*Univ. Bordeaux*)
- Charles Deledalle (*Brain Corp*)
- Hugo Larochelle (*Google Brain*)
- Andrew Ng (*Stanford*)
- Baptiste Pesquet (*ENSC*)

Objectives and organization

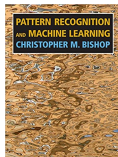
Objectives

- Understand the key elements of **machine learning theory**
- Know some machine learning **algorithms** (not all)
- Be able to **implement** and **use** them in different **applications**

Organization

- **24 hours** total (2h/3h sessions), mix of lectures and practices
- IS318 page on *Thor* for information and class materials
- **Grading**: ~20% assiduity, ~80% practices

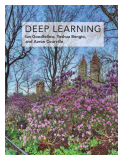
References



Pattern Recognition and Machine Learning

Christopher M. Bishop

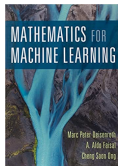
<https://www.microsoft.com/en-us/research/people/cmbishop/prml-book/>



Deep Learning

Ian Goodfellow, Yoshua Bengio, Aaron Courville

<https://www.deeplearningbook.org>



Mathematics for Machine Learning

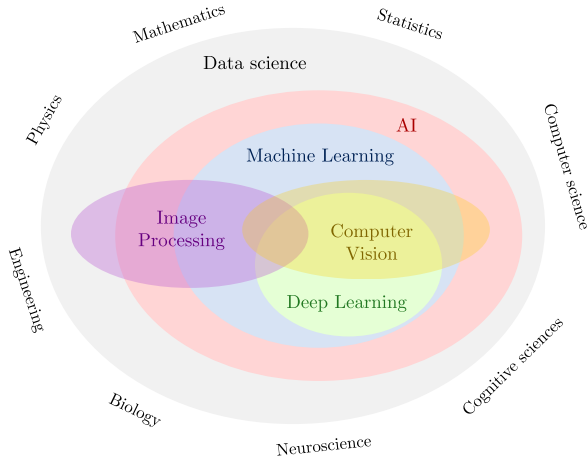
Marc Peter Deisenroth, A. Aldo Faisal, Cheng Soon Ong

<https://mml-book.com>

Introduction

What is machine learning?


A multidisciplinary subject



Definitions

“Learning is any process by which a system improves performance from experience.”

— **Herbert Simon**



Pavlov's dog
(Mark Stivers, 2003)

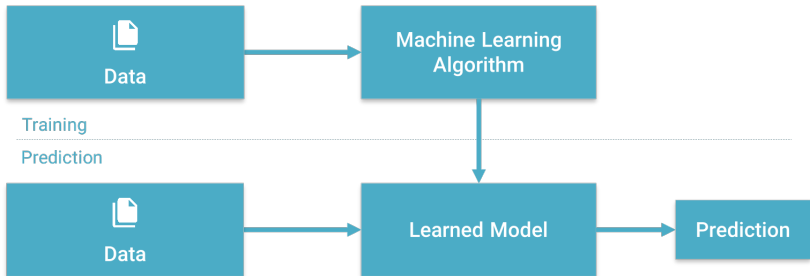
*“A computer program is said to learn from **experience E** with respect to some class of **tasks T** and performance **measure P** , if its performance at tasks in T , as measured by P , improves with experience E .”*

— **Tom Mitchell**

Learn from data

Machine learning provides **various techniques** that can learn from and make predictions on data.

Most of them follow the same general structure:



Main ingredients

To learn from examples, we will need:

- 1 Training data (examples):

$$\mathcal{D}_{\text{train}} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}, \mathbf{x}_i \in \mathbb{R}^d$$

- 2 Model (machine or program):

$$\underbrace{\mathbf{x}}_{\text{input data}} \rightarrow \underbrace{f(\mathbf{x}; \boldsymbol{\theta})}_{\text{function / algorithm}} \rightarrow \underbrace{\hat{\mathbf{y}}}_{\text{prediction}}$$

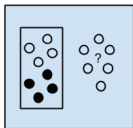
- 3 Loss, cost, objective function / energy:

$$\arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}; \{\mathbf{x}_1, \dots, \mathbf{x}_N\})$$

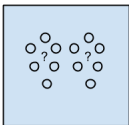
Terminology

- **Sample (Observation or Data):** item to process (e.g., classify).
Example: an individual, a document, a picture, a sound, a video, etc.
- **Features (Input):** set of distinct traits that can be used to describe each sample in a quantitative manner. Represented as a multi-dimensional vector usually denoted by x .
Example: size, weight, frequency, color, etc.
- **Training set:** set of data used to discover predictive relationships.
- **Testing set:** set of data used to assess the performance of a model.
- **Label or Prediction (Output):** the class or outcome assigned to a sample. The actual prediction is often denoted by \hat{y} and the ground truth value by y .
Example: cat/dog, temperature, price, etc.

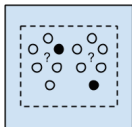
Learning approaches



Supervised Learning Algorithms



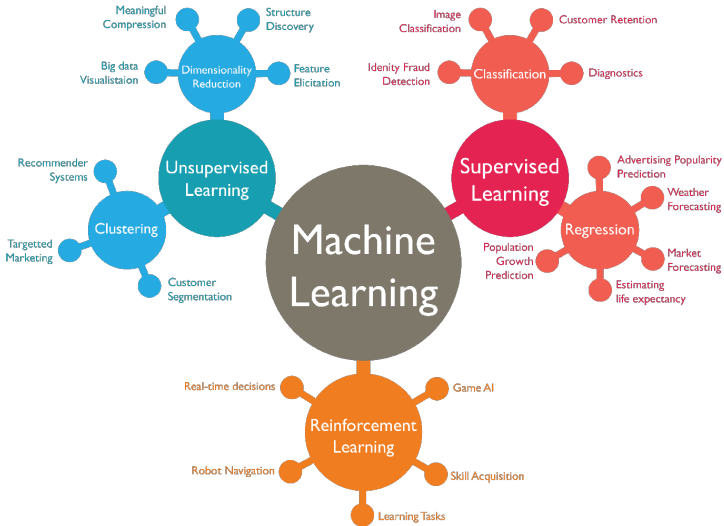
Unsupervised Learning Algorithms



Semi-supervised Learning Algorithms

- **Supervised learning:** learn from **labeled training data** $\mathcal{D}_{\text{train}} = \{(x_1, y_1), \dots, (x_N, y_N)\}$
- **Unsupervised learning:** discover patterns in **unlabeled training data** $\mathcal{D}_{\text{train}} = \{x_1, \dots, x_N\}$
- **Semi-supervised learning:** learn with a **small amount of labeled data and a large amount of unlabeled data**.
- **Reinforcement learning:** Learning based on **feedback** or reward \rightarrow IS320 class.

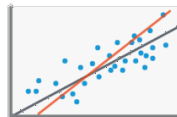
Different tasks



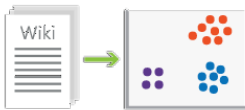
Problem types




Classification
(supervised – predictive)



Regression
(supervised – predictive)



Clustering
(unsupervised – descriptive)



Anomaly Detection
(unsupervised – descriptive)

No free lunch theorem

No free lunch theorem

- Generic theorem in optimization theory and machine learning
- It states that, *averaged over all possible problems and without any prior knowledge, there is no objectively better learning algorithm*
- For example, a very fancy learning algorithm will, on average, perform as well as taking random decisions



→ In practice, it means that we must design and apply solutions that are **biased towards the problem at hand**


- ① Fundamentals and linear regression
- ② Probabilistic modeling
- ③ Linear classifiers
- ④ Support vector machines
- ⑤ Decision trees
- ⑥ Combination of models

Fundamentals and linear regression

Example: Polynomial linear regression

Linear regression problem in 1D

- Input (feature): scalar $x \in \mathbb{R}$
- Output (target): scalar $y \in \mathbb{R}$



Train set: N labeled examples (supervised setting)


$$\mathcal{D}_{\text{train}} = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

Objective: learn a function f able to make prediction \hat{y} for input x

$$\hat{y} = f(x)$$

Model

Polynomial linear regression: we assume that good predictions follow a polynomial form



Our **model** f can be defined as:

$$f(x; \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Dx^D = \sum_{d=0}^D w_dx^d$$


where $\mathbf{w} = [w_0, \dots, w_D]^T \in \mathbb{R}^{D+1}$ are the **parameters** of the model.

Loss function

Learning problem:

How to find a “good” w ?

→ Find w^* minimizing the difference between (x_i, y_i) pairs in $\mathcal{D}_{\text{train}}$



Mean squared error: we are looking for $w^* = \arg \min_w \mathcal{L}(w)$ where

$$\mathcal{L}(w) = \frac{1}{N} \sum_{i=1}^N (f(x_i; w) - y_i)^2$$

This loss function is **convex**, so there is a single global minimum.

Optimization (1/2)

From the training set $\mathcal{D}_{\text{train}}$ we introduce:

- The data matrix $\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^D \\ \vdots & \vdots & \vdots & \dots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^D \end{bmatrix}$
- The target vector $\mathbf{y} = [y_1 \quad \dots \quad y_N]^\top$

Then, the mean squared error can be expressed as:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (f(x_i; \mathbf{w}) - y_i)^2 = \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

To minimize, because the function is convex, we can set the gradient to 0:

$$\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = 0 \Leftrightarrow \nabla_{\mathbf{w}} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = 0$$

Optimization (2/2)

We have¹:

$$\begin{aligned}\nabla_w \|Xw - y\|^2 = 0 &\Leftrightarrow \nabla_w (Xw - y)^\top (Xw - y) = 0 \\ &\Leftrightarrow \nabla_w w^\top X^\top Xw - 2w^\top X^\top y + y^\top y = 0 \\ &\Leftrightarrow 2X^\top Xw - 2X^\top y = 0 \\ &\Leftrightarrow X^\top Xw = X^\top y\end{aligned}$$

Assuming $X^\top X$ is invertible², we find:

$$w = (X^\top X)^{-1} X^\top y$$

This solution is called the **normal equation**.

¹For detailed derivations, see:

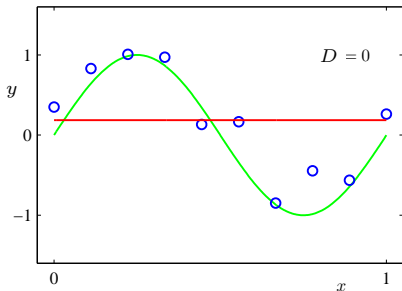
<https://eli.thegreenplace.net/2014/derivation-of-the-normal-equation-for-linear-regression>

²This is the case if we have $N \geq D$ distinct x_i

Underfitting and overfitting

How to choose the value of D ?

$$f(x; \mathbf{w}) = \sum_{d=0}^D w_d x^d$$



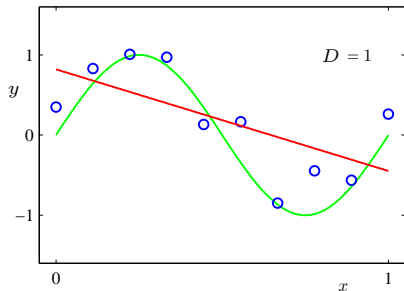
Too small: large error on the train set, too simple representation
→ **Underfitting**

D is called an **hyperparameter** of the model

Underfitting and overfitting

How to choose the value of D ?

$$f(x; \mathbf{w}) = \sum_{d=0}^D w_d x^d$$



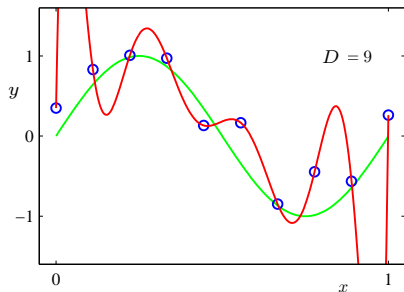
Too small: large error on the train set, too simple representation
→ **Underfitting**

D is called an **hyperparameter** of the model

Underfitting and overfitting

How to choose the value of D ?

$$f(x; \mathbf{w}) = \sum_{d=0}^D w_d x^d$$



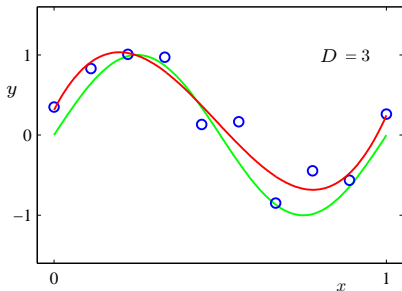
Too large: the model learns “by heart”, too complex representation
→ **Overfitting**

D is called an **hyperparameter** of the model

Underfitting and overfitting

How to choose the value of D ?

$$f(x; \mathbf{w}) = \sum_{d=0}^D w_d x^d$$



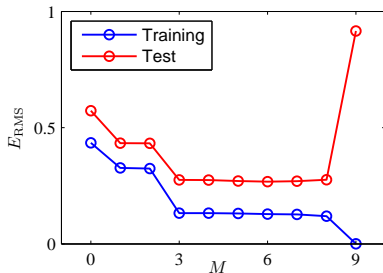
We want a good compromise that finds the general trend of the relationship, but without the noise, to **generalize** on new test data

D is called an **hyperparameter** of the model

Model capacity

Model capacity: ability of a model to learn the training set “by heart”

- Example: the larger D , the larger the capacity

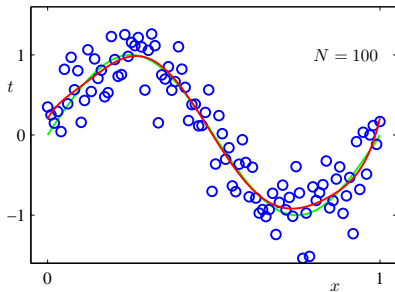
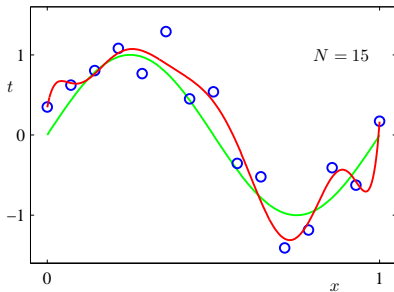


As the capacity increases, the difference between training error and test error is likely to increase

$$E_{RMS} = \sqrt{\mathcal{L}(\mathbf{w}^*)}$$

Generalization

As the **amount of training data** increases, the better the trained model will **generalize**.



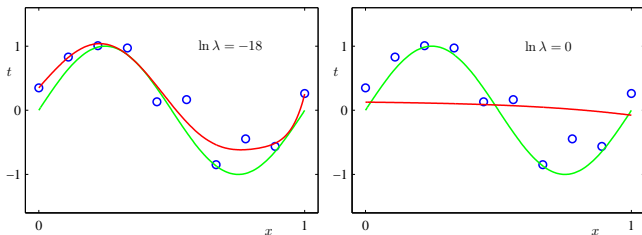
Regularization

Regularization: add a penalty term to the mean squared error

$$\tilde{\mathcal{L}}(\mathbf{w}) = \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2$$

with $\|\mathbf{w}\|^2 = \mathbf{w}^\top \mathbf{w} = w_0^2 + w_1^2 + \dots + w_N^2$ (*weight decay penalty*).

λ is another **hyperparameter**, like D , which allows to control the capacity of the linear regression model.



With weight decay, solving for \mathbf{w} gives: $\mathbf{w} = (\lambda \mathbf{I} + \mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$

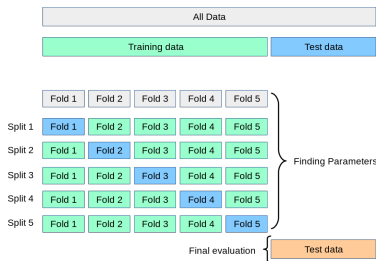
Model selection

How to choose the value of hyperparameters?

- Keep some training data aside in a **validation set** $\mathcal{D}_{\text{valid}}$
- For example: split data into 80% for $\mathcal{D}_{\text{train}}$ and 20% for $\mathcal{D}_{\text{valid}}$
- For different values of hyperparameters, choose the model that gives the best performance on $\mathcal{D}_{\text{valid}}$

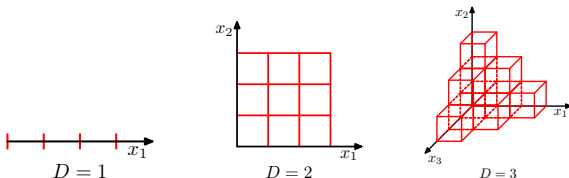
Cross validation

- Repeat the procedure for different $\mathcal{D}_{\text{train}}$ and $\mathcal{D}_{\text{valid}}$ splits
- Average the results of each split to choose hyperparameters
- Useful when training data is limited



Curse of dimensionality

- To generalize, we would like to have enough training examples to “fill” the underlying **feature space**
- However, this number grows **exponentially** with the dimension D
 - To have 10 samples in each dimension, it would require 10^D samples
- In machine learning, $D > 100$ or $D > 1000$ is very common



- But all is not lost! There is usually regularity/structure in the data
 - Dimensionality reduction
 - Feature engineering/learning